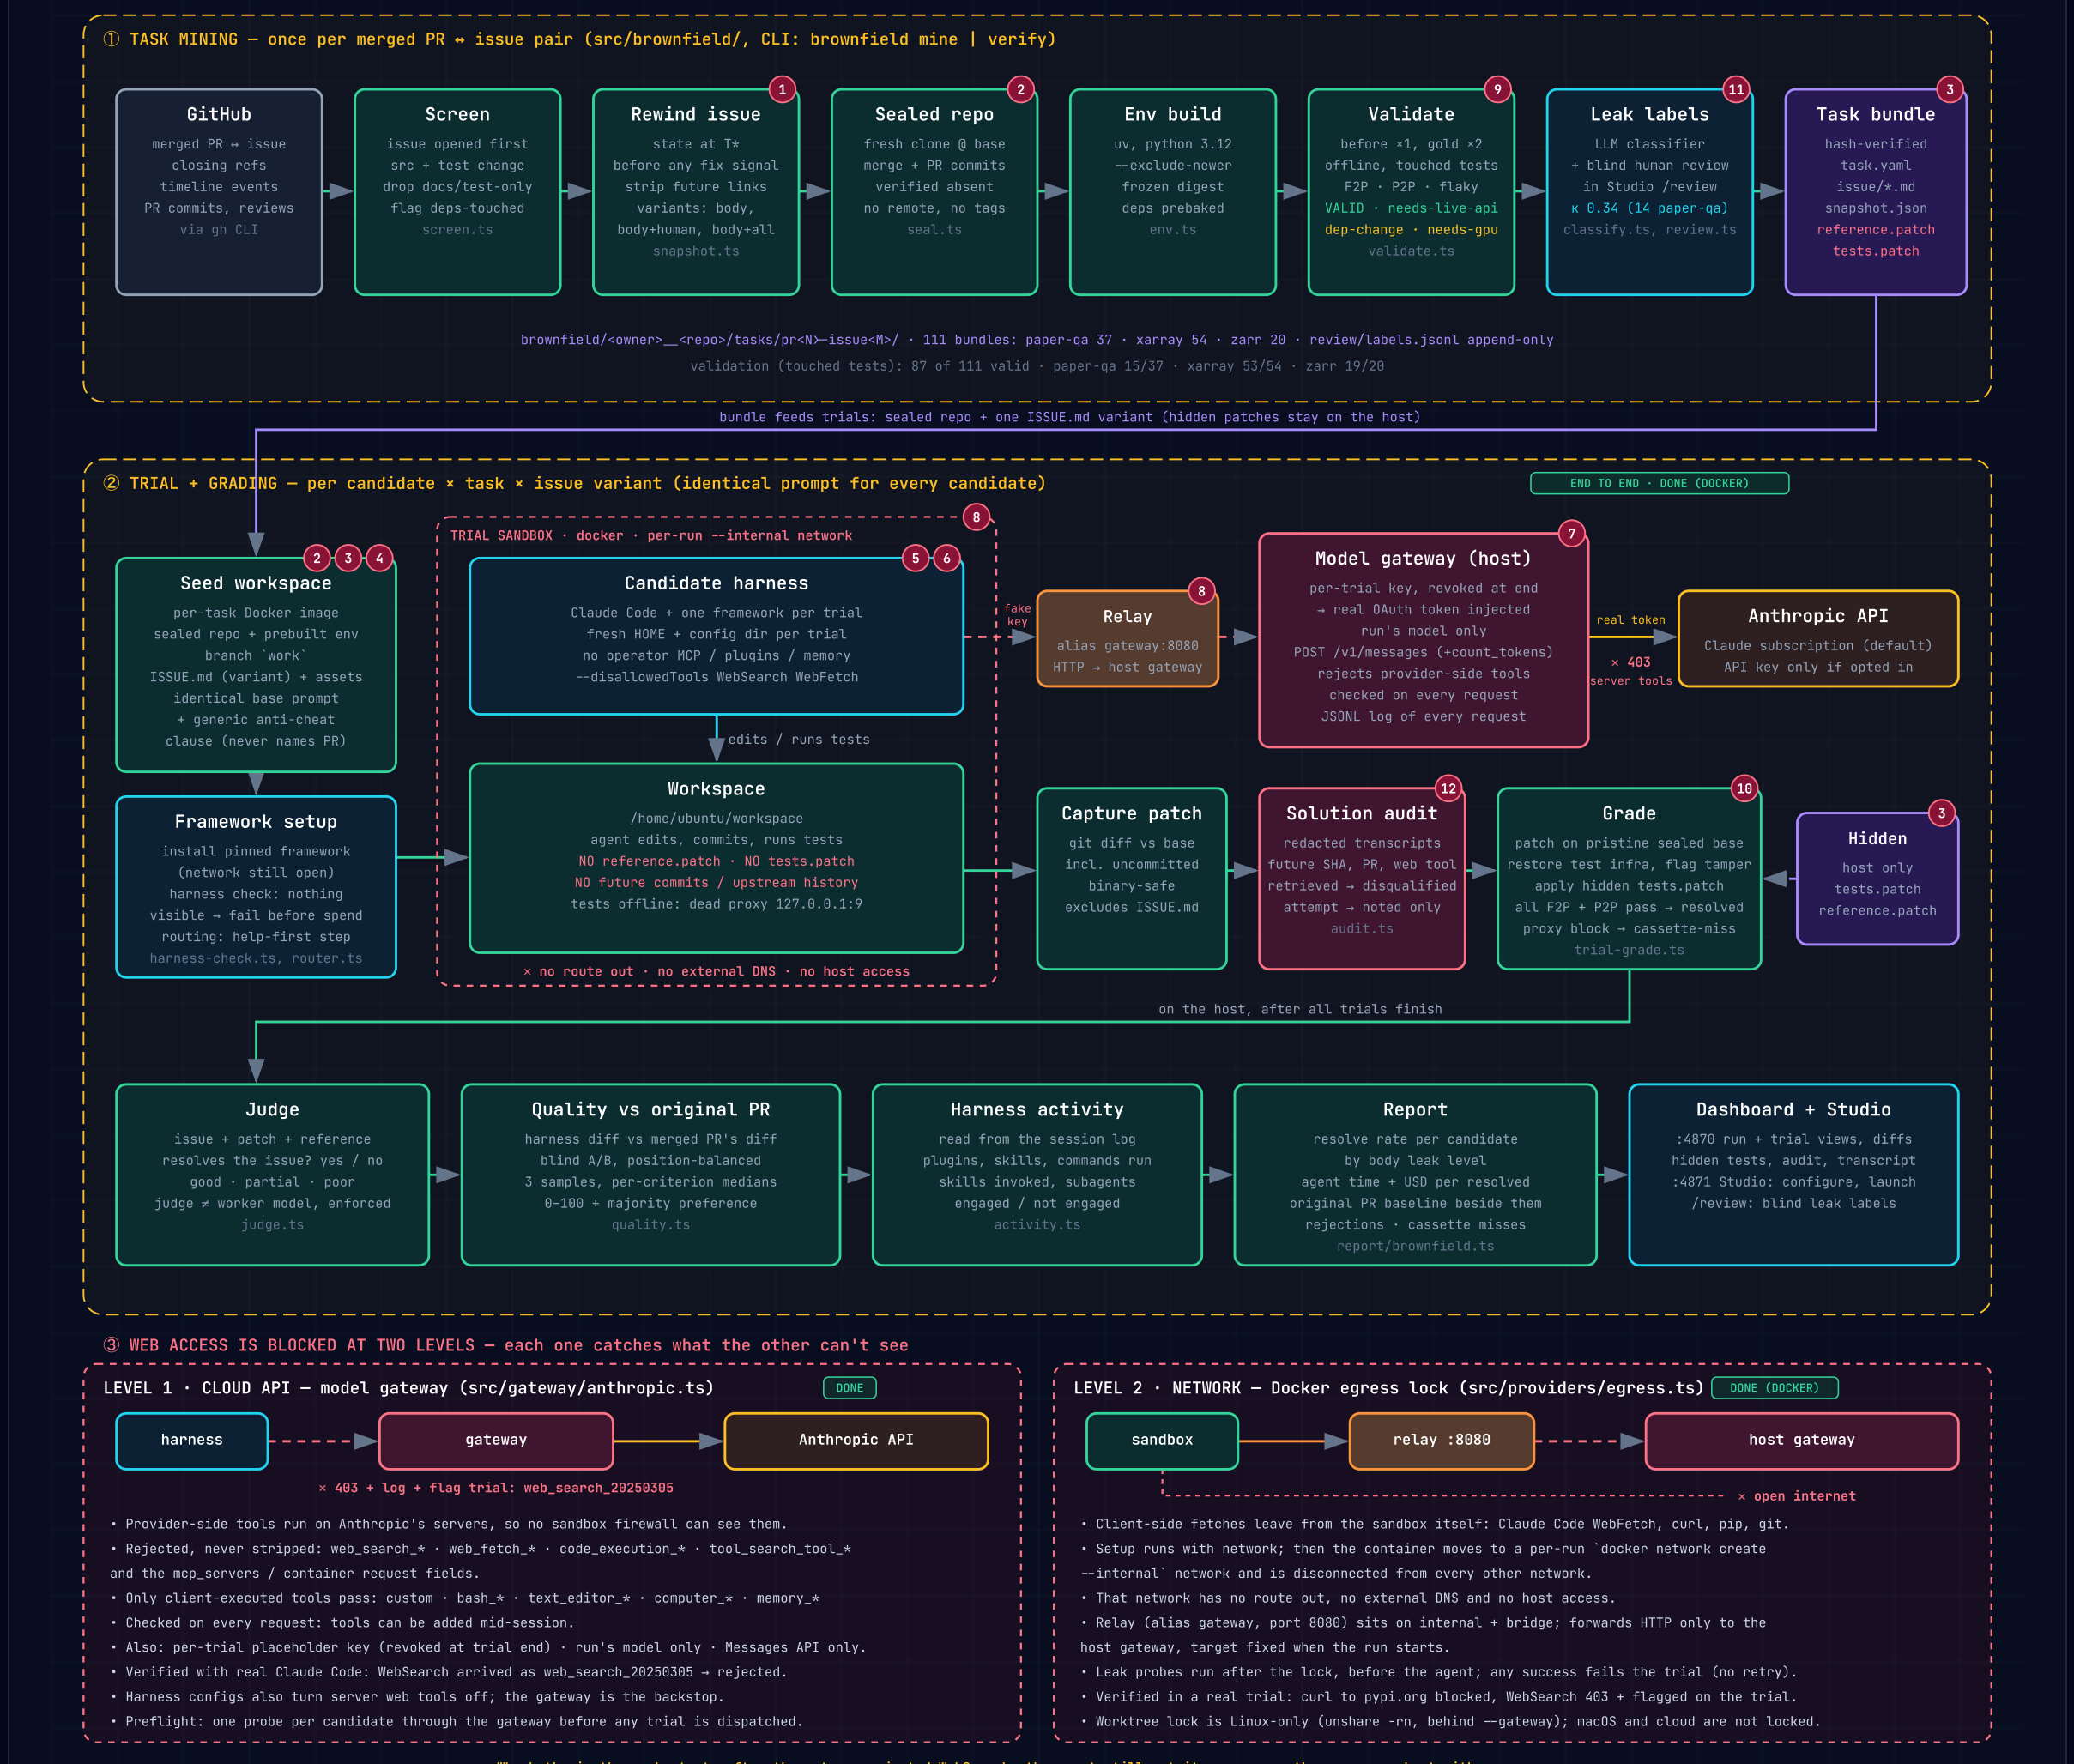
harness-eval · Brownfield Track

Testing coding-agent harnesses on real merged PRs: each issue is rewound to before the fix, solved in a sealed repo, and graded with the PR's hidden tests, without letting the agent see the answer.



How a task is made

- Take a merged PR that closes an issue opened before the fix work began
- Rewind the issue to T* and seal the repo at the PR's base commit
- Validate offline: the gold patch must turn hidden tests from failing to passing (F2P) without breaking others (P2P), twice
- Label how much each issue body/comment leaks the fix (LLM + blind human review)

How a harness is scored

- Every candidate gets the same sealed repo, ISSUE.md variant and base prompt; frameworks with no entry command get an identical help-first routing step
- Its diff vs base is captured; test files and pytest config are reset to base
- Hidden tests.patch applied; resolved only if every F2P and P2P test passes
- A transcript audit can disqualify a trial before grading; an independent judge and a blind A/B against the original PR add quality signals
- Results reported by body leak level, beside agent time, USD and the original PR baseline

Status (2026-11-01)

- Done: mining, touched-test validation, bundles, Studio review, pilot
- Done on Docker with Claude Code: run --brownfield end to end on the Claude subscription, egress lock, audit, grading, judge, quality vs original PR, reports, dashboard views
- Open: Codex through the gateway; worktree (macOS) and cloud egress; application egress + live-API task recovery; framework engagement audit (5.8); paired variant report
- Fixed: costs counted per step, timeouts reported as timeouts, engagement read from session evidence, judges see the whole patch
- Full-suite validation deferred to the cloud orchestrator